

## Target:

```
-- Get the set-up right
-- Set Transforms
-- Set Data Loader
-- Set Basic Working Code
-- Set Basic Training & Test Loop
-- Use batch normalisation
```

## Results:

```
-- Parameters: 6,383,818
-- Best Training Accuracy: 99.99
-- Best Test Accuracy: 99.61
```

## Analysis:

```
-- The accuracy is really good. There is no overfitting as the test accuracy
is increasing along with the training accuracy
-- model is really heavy. 6.3M parameters are really heavy
```

## ▼ Import libraries

```
from __future__ import print_function
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
```

## ▼ Data Transformations (without normalization)

```
# Train Phase transformations
train_transforms = transforms.Compose([
    transforms.ToTensor()
])

# Test Phase transformations
test_transforms = transforms.Compose([
```

```
transforms.ToTensor()  
1)
```

## Dataset and Creating Train/Test Split (without normalization)

```
train = datasets.MNIST('./data', train=True, download=True, transform=train_trans:  
test = datasets.MNIST('./data', train=False, download=True, transform=test_trans:
```

## Dataloader Arguments & Test/Train Dataloaders (without normalization)

```
SEED = 1  
  
# CUDA?  
cuda = torch.cuda.is_available()  
print("CUDA Available?", cuda)  
  
# For reproducibility  
torch.manual_seed(SEED)  
  
if cuda:
```

```

torch.cuda.manual_seed(SEED)

# dataloader arguments - something you'll fetch these from cmdprmt
dataloader_args = dict(shuffle=True, batch_size=128, num_workers=4, pin_memory=T:

# train dataloader
train_loader = torch.utils.data.DataLoader(train, **dataloader_args)

# test dataloader
test_loader = torch.utils.data.DataLoader(test, **dataloader_args)

CUDA Available? True
/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554:
warnings.warn(_create_warning_msg(

```

## ▼ Getting data statistics (without normalization)

We will use the mean and standard deviation that we get from code below to normalize the data

```

import numpy as np

train_data = train.train_data
train_data = train.transform(train_data.numpy())

print('[Train]')
print(' - Numpy Shape:', train.train_data.cpu().numpy().shape)
print(' - Tensor Shape:', train.train_data.size())
print(' - min:', torch.min(train_data))
print(' - max:', torch.max(train_data))
print(' - mean:', torch.mean(train_data))
print(' - std:', torch.std(train_data))
print(' - var:', torch.var(train_data))

dataiter = iter(train_loader)
images, labels = next(dataiter)

print(images.shape)
print(labels.shape)

# Let's visualize some of the images
%matplotlib inline
import matplotlib.pyplot as plt

plt.imshow(images[0].numpy().squeeze(), cmap='gray_r')

```

## ▼ Data Transformations (with normalization)

```
# Train Phase transformations
train_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

# Test Phase transformations
test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])
```

## ▼ Dataset and Creating Train/Test Split (with normalization)

```
train = datasets.MNIST('./data', train=True, download=True, transform=train_trans:
test = datasets.MNIST('./data', train=False, download=True, transform=test_trans:
```

## ▼ Dataloader Arguments & Test/Train Dataloaders (with normalization)

```
SEED = 1

# CUDA?
cuda = torch.cuda.is_available()
print("CUDA Available?", cuda)

# For reproducibility
torch.manual_seed(SEED)

if cuda:
    torch.cuda.manual_seed(SEED)

# dataloader arguments - something you'll fetch these from cmdprmt
dataloader_args = dict(shuffle=True, batch_size=128, num_workers=4, pin_memory=T:

# train dataloader
train_loader = torch.utils.data.DataLoader(train, **dataloader_args)

# test dataloader
test_loader = torch.utils.data.DataLoader(test, **dataloader_args)

CUDA Available? True
```

## ▼ Getting data statistics (with normalization)

We will use the mean and standard deviation that we get from code below to normalize the data

```
import numpy as np

train_data = train.train_data
train_data = train.transform(train_data.numpy())

print('[Train]')
print(' - Numpy Shape:', train.train_data.cpu().numpy().shape)
print(' - Tensor Shape:', train.train_data.size())
print(' - min:', torch.min(train_data))
print(' - max:', torch.max(train_data))
print(' - mean:', torch.mean(train_data))
print(' - std:', torch.std(train_data))
print(' - var:', torch.var(train_data))

dataiter = iter(train_loader)
```

```
images, labels = next(dataiter)

print(images.shape)
print(labels.shape)

# Let's visualize some of the images
%matplotlib inline
import matplotlib.pyplot as plt

plt.imshow(images[0].numpy().squeeze(), cmap='gray_r')
```

## Model

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()

        #input block
        self.convblock1 = nn.Sequential(nn.Conv2d(in_channels = 1, out_channels = 32,
                                                    kernel_size = 3, padding = 1),
                                         nn.BatchNorm2d(32),
                                         nn.ReLU()) #R_in = 1, C_in = 28, K = 3, P = 1

        #conv block 1
        self.convblock2 = nn.Sequential(nn.Conv2d(in_channels = 32, out_channels = 64,
                                                    kernel_size = 3, padding = 1),
                                         nn.BatchNorm2d(64),
                                         nn.ReLU()) #R_in = 3, C_in = 28, K = 3, P = 1
```

```

#conv block 2
self.convblock3 = nn.Sequential(nn.Conv2d(in_channels = 64, out_channels = 128,
                                           kernel_size=3, padding=1, bias=True),
                                nn.BatchNorm2d(128),
                                nn.ReLU()) #R_in = 5, C_in = 28, K = 3, P = 1

#transition block1
self.convblock4 = nn.Sequential(nn.Conv2d(in_channels = 128, out_channels = 256,
                                           kernel_size=3, padding=1, bias=True),
                                nn.BatchNorm2d(256),
                                nn.ReLU()) #R_in = 7, C_in = 28, K = 3, P = 1

self.pool1 = nn.MaxPool2d(2, 2) #R_in = 9, C_in = 28, K = 2, P = 0, S = 2, J_in = 2, J_out = 1

#conv block 3
self.convblock5 = nn.Sequential(nn.Conv2d(in_channels = 256, out_channels = 512,
                                           kernel_size=3, padding=1, bias=True),
                                nn.BatchNorm2d(512),
                                nn.ReLU()) #R_in = 9, C_in = 14, K = 3, P = 1

#conv block 4
self.convblock6 = nn.Sequential(nn.Conv2d(in_channels = 512, out_channels = 1024,
                                           kernel_size=3, padding=1, bias=True),
                                nn.BatchNorm2d(1024),
                                nn.ReLU()) #R_in = 13, C_in = 14, K = 3, P = 1

#gap layer
self.gap = nn.Sequential(
    nn.AvgPool2d(kernel_size=4)) #R_in = 17, C_in = 14, K = 4, P = 1, S = 1

#output block
self.convblock7 = nn.Sequential(nn.Conv2d(in_channels = 1024, out_channels = 10,
                                           kernel_size=3, padding=1, bias=True),
                                nn.ReLU()) #R_in = 23, C_in = 14, K = 3, P = 0, S = 1, J_in = 2, J_out = 10

def forward(self, x):
    x = self.convblock1(x)
    x = self.convblock2(x)
    x = self.convblock3(x)
    x = self.convblock4(x)
    x = self.pool1(x)
    x = self.convblock5(x)
    x = self.convblock6(x)
    x = self.gap(x)
    x = self.convblock7(x)
    x = x.view(-1, 10)
    return F.log_softmax(x, dim=-1)

```

## Model parameters

```

!pip install torchsummary
from torchsummary import summary

use_cuda = torch.cuda.is_available()
device = torch.device("cuda" if use_cuda else "cpu")

```

```

model = Net().to(device)
summary(model, input_size=(1, 28, 28))
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-research-artifacts/simple/
Requirement already satisfied: torchsummary in /usr/local/lib/python3.8/dist-packages (1.1.0)

```

| Layer (type)   | Output Shape       | Param #   |
|----------------|--------------------|-----------|
| Conv2d-1       | [-1, 32, 28, 28]   | 320       |
| BatchNorm2d-2  | [-1, 32, 28, 28]   | 64        |
| ReLU-3         | [-1, 32, 28, 28]   | 0         |
| Conv2d-4       | [-1, 64, 28, 28]   | 18,496    |
| BatchNorm2d-5  | [-1, 64, 28, 28]   | 128       |
| ReLU-6         | [-1, 64, 28, 28]   | 0         |
| Conv2d-7       | [-1, 128, 28, 28]  | 73,856    |
| BatchNorm2d-8  | [-1, 128, 28, 28]  | 256       |
| ReLU-9         | [-1, 128, 28, 28]  | 0         |
| Conv2d-10      | [-1, 256, 28, 28]  | 295,168   |
| BatchNorm2d-11 | [-1, 256, 28, 28]  | 512       |
| ReLU-12        | [-1, 256, 28, 28]  | 0         |
| MaxPool2d-13   | [-1, 256, 14, 14]  | 0         |
| Conv2d-14      | [-1, 512, 14, 14]  | 1,180,160 |
| BatchNorm2d-15 | [-1, 512, 14, 14]  | 1,024     |
| ReLU-16        | [-1, 512, 14, 14]  | 0         |
| Conv2d-17      | [-1, 1024, 14, 14] | 4,719,616 |
| BatchNorm2d-18 | [-1, 1024, 14, 14] | 2,048     |
| ReLU-19        | [-1, 1024, 14, 14] | 0         |
| AvgPool2d-20   | [-1, 1024, 3, 3]   | 0         |
| Conv2d-21      | [-1, 10, 1, 1]     | 92,170    |

```

=====
Total params: 6,383,818
Trainable params: 6,383,818
Non-trainable params: 0
=====
Input size (MB): 0.00
Forward/backward pass size (MB): 15.96
Params size (MB): 24.35
Estimated Total Size (MB): 40.31
=====

```

## Training and Testing

```

from tqdm import tqdm

train_losses = []
test_losses = []
train_acc = []
test_acc = []

def train(model, device, train_loader, optimizer, epoch):
    model.train()
    pbar = tqdm(train_loader)
    correct = 0
    processed = 0
    for batch_idx, (data, target) in enumerate(pbar):
        # get samples

```



```

data, target = data.to(device), target.to(device)

# Init
optimizer.zero_grad()
# In PyTorch, we need to set the gradients to zero before starting to do backpropagation
# Because of this, when you start your training loop, ideally you should zero the gradients again
# Predict
y_pred = model(data)

# Calculate loss
loss = F.nll_loss(y_pred, target)
train_losses.append(loss)

# Backpropagation
loss.backward()
optimizer.step()

# Update pbar-tqdm

pred = y_pred.argmax(dim=1, keepdim=True) # get the index of the max log-probability
correct += pred.eq(target.view_as(pred)).sum().item()
processed += len(data)

pbar.set_description(desc= f'Loss={loss.item():.4f} Batch_id={batch_idx} Accuracy={correct/processed:.2f}')
train_acc.append(100*correct/processed)

def test(model, device, test_loader):
    model.eval()
    test_loss = 0
    correct = 0
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            test_loss += F.nll_loss(output, target, reduction='sum').item() # sum of batch loss
            pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-probability
            correct += pred.eq(target.view_as(pred)).sum().item()

    test_loss /= len(test_loader.dataset)
    test_losses.append(test_loss)

    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.2f}%)\n'.format
          test_loss, correct, len(test_loader.dataset),
          100. * correct / len(test_loader.dataset))

    test_acc.append(100. * correct / len(test_loader.dataset))

from torch.optim.lr_scheduler import StepLR

model = Net().to(device)
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
# scheduler = StepLR(optimizer, step_size=6, gamma=0.1)

```

```

EPOCHS = 15
for epoch in range(EPOCHS):
    print("EPOCH:", epoch)
    train(model, device, train_loader, optimizer, epoch)
    # scheduler.step()
    test(model, device, test_loader)

    EPOCH: 0
    Loss=0.038374871015548706 Batch_id=468 Accuracy=96.38: 100%|██████████| 469

    Test set: Average loss: 0.0418, Accuracy: 9870/10000 (98.70%)

    EPOCH: 1
    Loss=0.02019202709197998 Batch_id=468 Accuracy=98.98: 100%|██████████| 469/

    Test set: Average loss: 0.0276, Accuracy: 9912/10000 (99.12%)

    EPOCH: 2
    Loss=0.005100666545331478 Batch_id=468 Accuracy=99.27: 100%|██████████| 469

    Test set: Average loss: 0.0336, Accuracy: 9898/10000 (98.98%)

    EPOCH: 3
    Loss=0.004298883955925703 Batch_id=468 Accuracy=99.38: 100%|██████████| 469

    Test set: Average loss: 0.0266, Accuracy: 9921/10000 (99.21%)

    EPOCH: 4
    Loss=0.0004116976633667946 Batch_id=468 Accuracy=99.53: 100%|██████████| 46

    Test set: Average loss: 0.0301, Accuracy: 9899/10000 (98.99%)

    EPOCH: 5
    Loss=0.014014105312526226 Batch_id=468 Accuracy=99.55: 100%|██████████| 469

    Test set: Average loss: 0.0203, Accuracy: 9938/10000 (99.38%)

    EPOCH: 6
    Loss=0.007350246887654066 Batch_id=468 Accuracy=99.63: 100%|██████████| 469

    Test set: Average loss: 0.0224, Accuracy: 9932/10000 (99.32%)

    EPOCH: 7
    Loss=0.0031526677776128054 Batch_id=468 Accuracy=99.73: 100%|██████████| 46

    Test set: Average loss: 0.0180, Accuracy: 9942/10000 (99.42%)

    EPOCH: 8
    Loss=0.003658372675999999 Batch_id=468 Accuracy=99.75: 100%|██████████| 469

    Test set: Average loss: 0.0239, Accuracy: 9923/10000 (99.23%)

    EPOCH: 9
    Loss=0.00021265355462674052 Batch_id=468 Accuracy=99.78: 100%|██████████| 4

    Test set: Average loss: 0.0188, Accuracy: 9945/10000 (99.45%)

```

EPOCH: 10

Loss=0.010211896151304245 Batch\_id=468 Accuracy=99.86: 100% |  | 469

Test set: Average loss: 0.0155, Accuracy: 9949/10000 (99.49%)

EPOCH: 11

Loss=0.0006745135760866106 Batch\_id=468 Accuracy=99.90: 100% |  | 46

