

## ▼ Import libraries

```
from __future__ import print_function
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms

from google.colab import drive
drive.mount('/content/drive', force_remount = True)

Mounted at /content/drive

%cd 'drive/MyDrive/EVA8/Session 5 - Assignment'

/content/drive/MyDrive/EVA8/Session 5 - Assignment

from model import NetWithGN, NetWithLN, NetWithBN
```

## ▼ Data Transformations (without normalization)

```
# Train Phase transformations
train_transforms = transforms.Compose([
    transforms.RandomRotation((-7.0, 7.0), fill=(1,)),
    transforms.ToTensor()
])

# Test Phase transformations
test_transforms = transforms.Compose([
    transforms.ToTensor()
])
```

A Google Drive error has occurred. ✕

## ▼ Dataset and Creating Train/Test Split (without normalization)

```
train = datasets.MNIST('./data', train=True, download=True, transform=train_transforms)
test = datasets.MNIST('./data', train=False, download=True, transform=test_transforms)
```

## ▼ Dataloader Arguments & Test/Train Dataloaders (without normalization)

```
SEED = 1

# CUDA?
cuda = torch.cuda.is_available()
print("CUDA Available?", cuda)

# For reproducibility
torch.manual_seed(SEED)

if cuda:
    torch.cuda.manual_seed(SEED)

# dataloader arguments - something you'll fetch these from cmdprmt
dataloader_args = dict(shuffle=True, batch_size=128, num_workers=4, pin_memory=True) if cuda else dict(shuffle=True, batch

# train dataloader
train_loader = torch.utils.data.DataLoader(train, **dataloader_args)

# test dataloader
test_loader = torch.utils.data.DataLoader(test, **dataloader_args)

CUDA Available? True
/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554: UserWarning: This DataLoader will create 4
warnings.warn(_create_warning_msg(
```

## ▼ Data Transformations (with normalization)

```
# Train Phase transformations
train_transforms = transforms.Compose([
    transforms.RandomRotation((-7.0, 7.0), fill=(1,)),
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

# Test Phase transformations
test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])
```

## ▼ Dataset and Creating Train/Test Split (with normalization)

```
train = datasets.MNIST('./data', train=True, download=True, transform=train_transforms)
test = datasets.MNIST('./data', train=False, download=True, transform=test_transforms)
```

## ▼ Dataloader Arguments & Test/Train Dataloaders (with normalization)

```
SEED = 1

# CUDA?
cuda = torch.cuda.is_available()
print("CUDA Available?", cuda)
```

A Google Drive error has occurred. ✕

```
if cuda:
    torch.cuda.manual_seed(SEED)

# dataloader arguments - something you'll fetch these from cmdprmt
dataloader_args = dict(shuffle=True, batch_size=128, num_workers=4, pin_memory=True) if cuda else dict(shuffle=True, batch

# train dataloader
train_loader = torch.utils.data.DataLoader(train, **dataloader_args)

# test dataloader
test_loader = torch.utils.data.DataLoader(test, **dataloader_args)

    CUDA Available? True
```

## ▼ plot some images to see which image augmentation to use (with normalization)

We will plot some images to see which image augmentation technique we can use

```
dataiter = iter(train_loader)
images, labels = next(dataiter)

# Let's visualize some of the images
%matplotlib inline
import matplotlib.pyplot as plt

figure = plt.figure()
num_of_images = 60
for index in range(1, num_of_images + 1):
    plt.subplot(6, 10, index)
    plt.axis('off')
    plt.imshow(images[index].numpy().squeeze(), cmap='gray_r')
```

### Training and Testing

```
from tqdm import tqdm
```

```
class TrainAndTest:
```

```
    def __init__(self):
```

```
        self.train_losses = []
```

```
        self.test_losses = []
```

```
        self.train_acc = []
```

```
        self.test_acc = []
```

```
        self.misclassified_images = []
```

```
    def train(self, model, device, train_loader, optimizer, epoch, task):
```

```
        """
```

```
        this function will train the model
```

```
        params model: model to be used for training
```

```
        params device: can be GPU or CPU
```

```
        params train_loader: training data
```

```
        optimiser: optimiser to be used
```

```
        epoch: epoch number
```

```
        task: task for which training will be performed (can be 'GN' for Group normalisation, 'LN' for Layer normalisation or
```

```
             If this is 'BN' then only we will do L1 regularisation)
```

```
        returns: train_acc, train_losses
```

```
        """
```

```
        model.train()
```

```
        pbar = tqdm(train_loader)
```

```
        correct = 0
```

A Google Drive error has occurred.



```
        in enumerate(pbar):
```

```
            # get samples
            data, target = data.to(device), target.to(device)
```

```
            # Init
```

```
            optimizer.zero_grad()
```

```
            # In PyTorch, we need to set the gradients to zero before starting to do backpropagation because PyTorch accumulate
```

```
            # Because of this, when you start your training loop, ideally you should zero out the gradients so that you do the p
```

```
            # Predict
```

```
            y_pred = model(data)
```

```
            # Calculate loss
```

```
            loss = F.nll_loss(y_pred, target)
```

```
            # Calculate 'L1' regularization loss if task == 'BN'
```

```
            if(task=='BN'):
```

```
                lambda_l1 = 0.01
```

```
                l1_reg = 0
```

```
                for param in model.parameters():
```

```
                    l1_reg += torch.norm(param, p=1)
```

```
                l1_loss = lambda_l1 * l1_reg
```

```
                loss += l1_loss
```

```
            self.train_losses.append(loss)
```

```
            # Backpropagation
```

```
            loss.backward()
```

```
            optimizer.step()
```

```
            # Update pbar-tqdm
```

```
            pred = y_pred.argmax(dim=1, keepdim=True) # get the index of the max log-probability
```

```
            correct += pred.eq(target.view_as(pred)).sum().item()
```

```
            processed += len(data)
```

```
            pbar.set_description(desc= f'Loss={loss.item()} Batch_id={batch_idx} Accuracy={100*correct/processed:0.2f}')
```

```
            self.train_acc.append(100*correct/processed)
```

```
        return self.train_acc, self.train_losses
```

```
    def test(self, model, device, test_loader):
```

```

"""
this function will test the model in test data
params model: model to be tested
params device: can be GPU or CPU
params test_loader: test data
returns: test_acc, test_losses and misclassified_images
"""

model.eval()
test_loss = 0
correct = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        output = model(data)
        test_loss += F.nll_loss(output, target, reduction='sum').item() # sum up batch loss
        pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-probability
        indices = [item[0] for item in (pred.eq(target.view_as(pred)) == False).nonzero().tolist()]
        mis_imgs = data[indices, :, :, :]
        self.misclassified_images.extend(mis_imgs)
        correct += pred.eq(target.view_as(pred)).sum().item()

test_loss /= len(test_loader.dataset)
self.test_losses.append(test_loss)

print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.2f}%)\n'.format(
    test_loss, correct, len(test_loader.dataset),
    100. * correct / len(test_loader.dataset)))

self.test_acc.append(100. * correct / len(test_loader.dataset))
return self.test_acc, self.test_losses, self.misclassified_images

use_cuda = torch.cuda.is_available()
device = torch.device("cuda" if use_cuda else "cpu")

def get_model(task_name):
    """
    this function will return the model based on the task_name
    """
    'BN'

    if (task_name=='GN'):
        print("You have chosen GN, Now we will perform Group Normalisation")
        model = NetWithGN().to(device)
    elif(task_name=='LN'):
        print("You have chosen LN, Now we will perform Layer Normalisation")
        model = NetWithLN().to(device)
    elif(task_name=='BN'):
        print("You have chosen BN, Now we will perform Batch Normalisation with L1 regularisation")
        model = NetWithBN().to(device)
    return model

from torch.optim.lr_scheduler import StepLR
def perform_tasks(task_name):
    """
    performs the train and test on the data
    params task_name: (GN, LN, BN)
    returns: train_acc, train_losses, test_acc, test_losses, misclassified_images
    """
    model = get_model(task_name) #get the model based on the task name (GN, LN, BN)
    optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9) #declare optimiser
    scheduler = StepLR(optimizer, step_size=6, gamma=0.1) #learning rate scheduler
    train_test_obj = TrainAndTest() #create the object of TrainAndTest class
    EPOCHS = 20
    for epoch in range(EPOCHS): #for each epoch
        print("EPOCH:", epoch)
        train_acc, train_losses = train_test_obj.train(model, device, train_loader, optimizer, epoch, task_name) #model train
        scheduler.step() #update weights
        test_acc, test_losses, misclassified_images = train_test_obj.test(model, device, test_loader) #test model using test
    return train_acc, train_losses, test_acc, test_losses, misclassified_images

import random
def show_misclassified_images(misclassified_images, task_name):
    """
    show 10 misclassified images
    """
    misclassified_images = random.sample(misclassified_images, 10)
    fig, axs = plt.subplots(nrows=5, ncols=2, figsize=(8, 12))
    axs = axs.ravel()

```

```

for i, ax in enumerate(axes):
    if i < 10:
        ax.imshow(misclassified_images[i].cpu().numpy().squeeze(), cmap='gray')
        ax.set_title('Image ' + str(i))

plt.tight_layout()
plt.suptitle('Misclassified images using Group Normalisation', y=1.02)
plt.show()

```

## Group Normalisation

```

task_name = 'GN'
train_acc_GN, train_losses_GN, test_acc_GN, test_losses_GN, misclassified_images_GN = perform_tasks(task_name)
train_losses_GN = [i.item() for i in train_losses_GN]

```

Test set: Average loss: 0.0212, Accuracy: 9921/10000 (99.21%)

EPOCH: 9

Loss=0.05961877107620239 Batch\_id=468 Accuracy=97.69: 100% ██████████ 469/469 [00:18<00:00, 25.92it/s]

Test set: Average loss: 0.0198, Accuracy: 9921/10000 (99.21%)

EPOCH: 10

Loss=0.04525679349899292 Batch\_id=468 Accuracy=97.66: 100% ██████████ 469/469 [00:18<00:00, 25.96it/s]

Test set: Average loss: 0.0205, Accuracy: 9927/10000 (99.27%)

EPOCH: 11

Loss=0.04970937967300415 Batch\_id=468 Accuracy=97.52: 100% ██████████ 469/469 [00:18<00:00, 25.47it/s]

Test set: Average loss: 0.0187, Accuracy: 9922/10000 (99.22%)

EPOCH: 12

Loss=0.06041562557220459 Batch\_id=468 Accuracy=97.67: 100% ██████████ 469/469 [00:18<00:00, 25.88it/s]

Test set: Average loss: 0.0187, Accuracy: 9928/10000 (99.28%)

A Google Drive error has occurred. ✕

Batch\_id=468 Accuracy=97.80: 100% ██████████ 469/469 [00:17<00:00, 26.20it/s]

Test set: Average loss: 0.0186, Accuracy: 9930/10000 (99.30%)

EPOCH: 14

Loss=0.08620817214250565 Batch\_id=468 Accuracy=97.74: 100% ██████████ 469/469 [00:18<00:00, 25.99it/s]

Test set: Average loss: 0.0188, Accuracy: 9928/10000 (99.28%)

EPOCH: 15

Loss=0.05799892917275429 Batch\_id=468 Accuracy=97.73: 100% ██████████ 469/469 [00:18<00:00, 26.00it/s]

Test set: Average loss: 0.0187, Accuracy: 9925/10000 (99.25%)

EPOCH: 16

Loss=0.06123394891619682 Batch\_id=468 Accuracy=97.65: 100% ██████████ 469/469 [00:18<00:00, 26.02it/s]

Test set: Average loss: 0.0186, Accuracy: 9926/10000 (99.26%)

EPOCH: 17

Loss=0.04270932078361511 Batch\_id=468 Accuracy=97.71: 100% ██████████ 469/469 [00:18<00:00, 25.69it/s]

Test set: Average loss: 0.0187, Accuracy: 9926/10000 (99.26%)

EPOCH: 18

Loss=0.06956575065851212 Batch\_id=468 Accuracy=97.74: 100% ██████████ 469/469 [00:17<00:00, 26.15it/s]

Test set: Average loss: 0.0186, Accuracy: 9926/10000 (99.26%)

EPOCH: 19

Loss=0.07965157181024551 Batch\_id=468 Accuracy=97.72: 100% ██████████ 469/469 [00:17<00:00, 26.13it/s]


Test set: Average loss: 0.0186, Accuracy: 9926/10000 (99.26%)

```

show_misclassified_images(misclassified_images_GN, 'Group Normalisation')

```

Misclassified images using Group Normalisation



A Google Drive error has occurred. ✕

```
task_name = 'LN'
train_acc_LN, train_losses_LN, test_acc_LN, test_losses_LN, misclassified_images_LN = perform_tasks(task_name)
train_losses_LN = [i.item() for i in train_losses_LN]
```

```

EPOCH: 16
Loss=0.10061508417129517 Batch_id=468 Accuracy=97.04: 100%|██████████| 469/469 [00:18<00:00, 25.05it/s]

Test set: Average loss: 0.0452, Accuracy: 9871/10000 (98.71%)

EPOCH: 17
Loss=0.09726373106241226 Batch_id=468 Accuracy=97.00: 100%|██████████| 469/469 [00:18<00:00, 25.14it/s]

Test set: Average loss: 0.0436, Accuracy: 9873/10000 (98.73%)

EPOCH: 18
Loss=0.11846256256103516 Batch_id=468 Accuracy=96.96: 100%|██████████| 469/469 [00:18<00:00, 25.10it/s]

Test set: Average loss: 0.0441, Accuracy: 9861/10000 (98.61%)

EPOCH: 19
Loss=0.07192239910364151 Batch_id=468 Accuracy=97.03: 100%|██████████| 469/469 [00:18<00:00, 25.40it/s]

Test set: Average loss: 0.0456, Accuracy: 9869/10000 (98.69%)

```

```
show_misclassified_images(misclassified_images_LN, 'Layer Normalisation')
```

Misclassified images using Group Normalisation



### Bacth Normalisation with L1 regularisation

```

task_name = 'BN'
train_acc_BN, train_losses_BN, test_acc_BN, test_losses_BN, misclassified_images_BN = perform_tasks(task_name)
train_losses_BN = [i.item() for i in train_losses_BN]

```

```
EPOCH: 9
Loss=0.7655919790267944 Batch_id=468 Accuracy=91.95: 100%|██████████| 469/469 [00:20<00:00, 23.42it/s]

Test set: Average loss: 0.3678, Accuracy: 8989/10000 (89.89%)

EPOCH: 10
Loss=0.8030099868774414 Batch_id=468 Accuracy=91.96: 100%|██████████| 469/469 [00:19<00:00, 23.83it/s]

Test set: Average loss: 0.1611, Accuracy: 9645/10000 (96.45%)

EPOCH: 11
Loss=0.7599208354949951 Batch_id=468 Accuracy=91.95: 100%|██████████| 469/469 [00:19<00:00, 23.93it/s]

Test set: Average loss: 0.1932, Accuracy: 9636/10000 (96.36%)

EPOCH: 12
Loss=0.6952192187309265 Batch_id=468 Accuracy=93.25: 100%|██████████| 469/469 [00:19<00:00, 23.78it/s]

Test set: Average loss: 0.1414, Accuracy: 9698/10000 (96.98%)

EPOCH: 13
Loss=0.7075787782669067 Batch_id=468 Accuracy=93.10: 100%|██████████| 469/469 [00:19<00:00, 23.49it/s]

Test set: Average loss: 0.1280, Accuracy: 9736/10000 (97.36%)

EPOCH: 14
Loss=0.6282301545143127 Batch_id=468 Accuracy=93.18: 100%|██████████| 469/469 [00:20<00:00, 23.43it/s]

Test set: Average loss: 0.1265, Accuracy: 9710/10000 (97.10%)

EPOCH: 15
Loss=0.6804661750793457 Batch_id=468 Accuracy=93.08: 100%|██████████| 469/469 [00:19<00:00, 23.73it/s]

Test set: Average loss: 0.1502, Accuracy: 9683/10000 (96.83%)

EPOCH: 16
Loss=0.6936613321304321 Batch_id=468 Accuracy=93.09: 100%|██████████| 469/469 [00:20<00:00, 23.33it/s]

Test set: Average loss: 0.1341, Accuracy: 9741/10000 (97.41%)

EPOCH: 17
Loss=0.744918704032898 Batch_id=468 Accuracy=93.08: 100%|██████████| 469/469 [00:20<00:00, 23.21it/s]

Test set: Average loss: 0.1376, Accuracy: 9713/10000 (97.13%)

EPOCH: 18
Loss=0.6583273410797119 Batch_id=468 Accuracy=93.43: 100%|██████████| 469/469 [00:19<00:00, 23.50it/s]

Test set: Average loss: 0.1245, Accuracy: 9744/10000 (97.44%)

EPOCH: 19
Loss=0.6359788179397583 Batch_id=468 Accuracy=93.46: 100%|██████████| 469/469 [00:19<00:00, 24.00it/s]

Test set: Average loss: 0.1252, Accuracy: 9746/10000 (97.46%)
```


```
show_misclassified_images(misclassified_images_BN, 'Batch Normalisation with L1 regularisation')
```

A Google Drive error has occurred.





Misclassified images using Group Normalisation



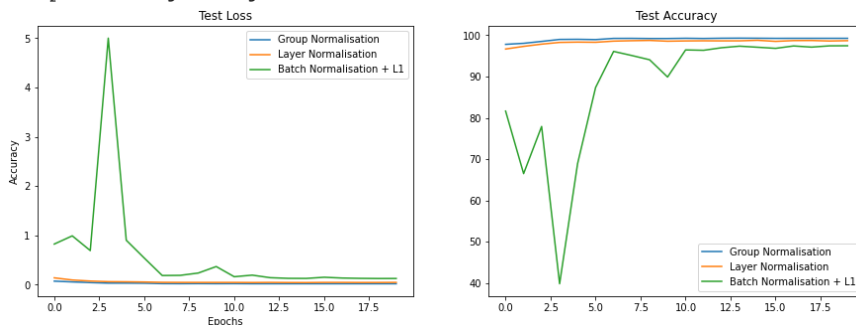
### Comparing losses and accuracies

```
fig, axs = plt.subplots(1,2,figsize=(15,5))
axs[0].plot(test_losses_GN, label='Group Normalisation')
axs[0].plot(test_losses_LN, label = 'Layer Normalisation')
axs[0].plot(test_losses_BN, label = 'Batch Normalisation + L1')
axs[0].set_title("Test Loss")
axs[0].set_xlabel('Epochs')
axs[0].set_ylabel('Loss')
axs[0].legend()

axs[1].plot(test_acc_GN, label='Group Normalisation')
axs[1].plot(test_acc_LN, label = 'Layer Normalisation')
axs[1].plot(test_acc_BN, label = 'Batch Normalisation + L1')
axs[1].set_title("Test Accuracy")
axs[1].set_xlabel('Epochs')
```

A Google Drive error has occurred. ✕

↳ <matplotlib.legend.Legend at 0x7f769e06ff40>



✓ 0s completed at 23:07

● ✕

A Google Drive error has occurred. ✕