

Target:

```
-- Get the set-up right
-- Set Transforms
-- Set Data Loader
-- Set Basic Working Code
-- Set Basic Training & Test Loop
-- Use batch normalisation
```

Results:

```
-- Parameters: 6,383,818
-- Best Training Accuracy: 99.97
-- Best Test Accuracy: 99.56
```

Analysis:

```
-- The accuracy is really good.
-- The model is starting to overfitting in last few epochs as the test accuracy is decreasing along with the training accuracy
-- model is really heavy. 6.3M parameters are really heavy
```

▾ Import libraries

```
from __future__ import print_function
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
```

▾ Data Transformations (without normalization)

```
# Train Phase transformations
train_transforms = transforms.Compose([
    transforms.ToTensor()
])

# Test Phase transformations
test_transforms = transforms.Compose([
    transforms.ToTensor()
])
```

▾ Dataset and Creating Train/Test Split (without normalization)

```
train = datasets.MNIST('./data', train=True, download=True, transform=train_transforms)
test = datasets.MNIST('./data', train=False, download=True, transform=test_transforms)
```

```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz
100% 9912422/9912422 [00:00<00:00, 18324015.82it/s]
Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz

```

▼ Dataloader Arguments & Test/Train Dataloaders (without normalization)

```

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
SEED = 1

# CUDA?
cuda = torch.cuda.is_available()
print("CUDA Available?", cuda)

# For reproducibility
torch.manual_seed(SEED)

if cuda:
    torch.cuda.manual_seed(SEED)

# dataloader arguments - something you'll fetch these from cmdprmt
dataloader_args = dict(shuffle=True, batch_size=128, num_workers=4, pin_memory=True) if cuda else dict(shuffle=True, batch

# train dataloader
train_loader = torch.utils.data.DataLoader(train, **dataloader_args)

# test dataloader
test_loader = torch.utils.data.DataLoader(test, **dataloader_args)

CUDA Available? True
/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554: UserWarning: This DataLoader will create 4
warnings.warn(_create_warning_msg(

```

▼ Getting data statistics (without normalization)

We will use the mean and standard deviation that we get from code below to normalize the data

```

import numpy as np

train_data = train.train_data
train_data = train.transform(train_data.numpy())

print('[Train]')
print(' - Numpy Shape:', train.train_data.cpu().numpy().shape)
print(' - Tensor Shape:', train.train_data.size())
print(' - min:', torch.min(train_data))
print(' - max:', torch.max(train_data))
print(' - mean:', torch.mean(train_data))
print(' - std:', torch.std(train_data))
print(' - var:', torch.var(train_data))

dataiter = iter(train_loader)
images, labels = next(dataiter)

print(images.shape)
print(labels.shape)

# Let's visualize some of the images
%matplotlib inline
import matplotlib.pyplot as plt


plt.imshow(images[0].numpy().squeeze(), cmap='gray_r')

```

```

/usr/local/lib/python3.8/dist-packages/torchvision/datasets/mnist.py:75: UserWarning: train_data has been renamed dat
warnings.warn("train_data has been renamed data")
[Train]
- Numpy Shape: (60000, 28, 28)
- Tensor Shape: torch.Size([60000, 28, 28])
- min: tensor(0.)
- max: tensor(1.)
- mean: tensor(0.1307)
- std: tensor(0.3081)
- var: tensor(0.0949)
torch.Size([128, 1, 28, 28])
torch.Size([128])
<matplotlib.image.AxesImage at 0x7f4d91edec70>

```



▼ Data Transformations (with normalization)

```

|      |      |      |
# Train Phase transformations
train_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

# Test Phase transformations
test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

```

▼ Dataset and Creating Train/Test Split (with normalization)

```

train = datasets.MNIST('./data', train=True, download=True, transform=train_transforms)
test = datasets.MNIST('./data', train=False, download=True, transform=test_transforms)

```

▼ Dataloader Arguments & Test/Train Dataloaders (with normalization)

```

SEED = 1

# CUDA?
cuda = torch.cuda.is_available()
print("CUDA Available?", cuda)

# For reproducibility
torch.manual_seed(SEED)

if cuda:
    torch.cuda.manual_seed(SEED)

# dataloader arguments - something you'll fetch these from cmdprmt
dataloader_args = dict(shuffle=True, batch_size=128, num_workers=4, pin_memory=True) if cuda else dict(shuffle=True, batch

# train dataloader
train_loader = torch.utils.data.DataLoader(train, **dataloader_args)

# test dataloader
test_loader = torch.utils.data.DataLoader(test, **dataloader_args)

CUDA Available? True

```

▼ Getting data statistics (with normalization)

We will use the mean and standard deviation that we get from code below to normalize the data

```

import numpy as np

train_data = train.train_data

```

```
train_data = train.transform(train_data.numpy())

print('[Train]')
print(' - Numpy Shape:', train.train_data.cpu().numpy().shape)
print(' - Tensor Shape:', train.train_data.size())
print(' - min:', torch.min(train_data))
print(' - max:', torch.max(train_data))
print(' - mean:', torch.mean(train_data))
print(' - std:', torch.std(train_data))
print(' - var:', torch.var(train_data))
```


```
dataiter = iter(train_loader)
images, labels = next(dataiter)
```

```
print(images.shape)
print(labels.shape)
```

```
# Let's visualize some of the images
%matplotlib inline
import matplotlib.pyplot as plt
```

```
plt.imshow(images[0].numpy().squeeze(), cmap='gray_r')
```

```
[Train]
- Numpy Shape: (60000, 28, 28)
- Tensor Shape: torch.Size([60000, 28, 28])
- min: tensor(-0.4242)
- max: tensor(2.8215)
- mean: tensor(-0.0001)
- std: tensor(1.0000)
- var: tensor(1.0001)
torch.Size([128, 1, 28, 28])
torch.Size([128])
<matplotlib.image.AxesImage at 0x7f4d910474f0>
```



Model

```
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()

        #input block
        self.convblock1 = nn.Sequential(nn.Conv2d(in_channels = 1, out_channels = 32, kernel_size = 3, padding = 1),
                                         nn.BatchNorm2d(32),
                                         nn.ReLU()) #R_in = 1, C_in = 28, K = 3, P = 1, S = 1, J_in = 1, J_out = 1, R_out = R_i

        #conv block 1
        self.convblock2 = nn.Sequential(nn.Conv2d(in_channels = 32, out_channels = 64, kernel_size = 3, padding = 1),
                                         nn.BatchNorm2d(64),
                                         nn.ReLU()) #R_in = 3, C_in = 28, K = 3, P = 1, S = 1, J_in = 1, J_out = 1, R_out = R_i

        #conv block 2
        self.convblock3 = nn.Sequential(nn.Conv2d(in_channels = 64, out_channels = 128, kernel_size = 3, padding = 1),
                                         nn.BatchNorm2d(128),
                                         nn.ReLU()) #R_in = 5, C_in = 28, K = 3, P = 1, S = 1, J_in = 1, J_out = 1, R_out = R_i

        #transition block1
        self.convblock4 = nn.Sequential(nn.Conv2d(in_channels = 128, out_channels = 256, kernel_size = 3, padding = 1),
                                         nn.BatchNorm2d(256),
                                         nn.ReLU()) #R_in = 7, C_in = 28, K = 3, P = 1, S = 1, J_in = 1, J_out = 1, R_out = R_i

        self.pool1 = nn.MaxPool2d(2, 2) #R_in = 9, C_in = 28, K = 2, P = 0, S = 2, J_in = 1, J_out = 2, R_out = R_in + (K-1)*J

        #conv block 3
        self.convblock5 = nn.Sequential(nn.Conv2d(in_channels = 256, out_channels = 512, kernel_size = 3, padding = 1),
                                         nn.BatchNorm2d(512),
                                         nn.ReLU()) #R_in = 10, C_in = 14, K = 3, P = 1, S = 1, J_in = 2, J_out = 2, R_out = R_i
```

```

#conv block 4
self.convblock6 = nn.Sequential(nn.Conv2d(in_channels = 512, out_channels = 1024, kernel_size = 3, padding = 1),
                                nn.BatchNorm2d(1024),
                                nn.ReLU()) #R_in = 14, C_in = 14, K = 3, P = 1, S = 1, J_in = 2, J_out = 2, R_out = R_

#gap layer
self.gap = nn.Sequential(
    nn.AvgPool2d(kernel_size=4)) #R_in = 18, C_in = 14, K = 4, P = 1, S = 1, J_in = 2, J_out = 2, R_out = R_in + (

#output block
self.convblock7 = nn.Sequential(nn.Conv2d(in_channels = 1024, out_channels = 10, kernel_size = 3, padding = 0))
                                #R_in = 24, C_in = 14, K = 3, P = 0, S = 1, J_in = 2, J_out = 2, R_out = R_in + (K-1)*J_in = 24+

def forward(self, x):
    x = self.convblock1(x)
    x = self.convblock2(x)
    x = self.convblock3(x)
    x = self.convblock4(x)
    x = self.pool1(x)
    x = self.convblock5(x)
    x = self.convblock6(x)
    x = self.gap(x)
    x = self.convblock7(x)
    x = x.view(-1, 10)
    return F.log_softmax(x, dim=-1)

```

Model parameters

```

!pip install torchsummary
from torchsummary import summary

```

```

use_cuda = torch.cuda.is_available()
device = torch.device("cuda" if use_cuda else "cpu")

```

```

model = Net().to(device)
summary(model, input_size = (1, 28, 28))

```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
Requirement already satisfied: torchsummary in /usr/local/lib/python3.8/dist-packages (1.5.1)

```

-----
Layer (type)          Output Shape          Param #
-----
Conv2d-1              [-1, 32, 28, 28]      320
BatchNorm2d-2         [-1, 32, 28, 28]      64
ReLU-3                [-1, 32, 28, 28]      0
Conv2d-4              [-1, 64, 28, 28]      18,496
BatchNorm2d-5         [-1, 64, 28, 28]      128
ReLU-6                [-1, 64, 28, 28]      0
Conv2d-7              [-1, 128, 28, 28]     73,856
BatchNorm2d-8         [-1, 128, 28, 28]     256
ReLU-9                [-1, 128, 28, 28]      0
Conv2d-10             [-1, 256, 28, 28]     295,168
BatchNorm2d-11        [-1, 256, 28, 28]     512
ReLU-12               [-1, 256, 28, 28]      0
MaxPool2d-13          [-1, 256, 14, 14]      0
Conv2d-14             [-1, 512, 14, 14]     1,180,160
BatchNorm2d-15        [-1, 512, 14, 14]     1,024
ReLU-16               [-1, 512, 14, 14]      0
Conv2d-17             [-1, 1024, 14, 14]    4,719,616
BatchNorm2d-18        [-1, 1024, 14, 14]     2,048
ReLU-19               [-1, 1024, 14, 14]      0
AvgPool2d-20          [-1, 1024, 3, 3]      0
Conv2d-21             [-1, 10, 1, 1]        92,170
-----
Total params: 6,383,818
Trainable params: 6,383,818
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 15.96
Params size (MB): 24.35
Estimated Total Size (MB): 40.31
-----

```

Training and Testing

```

from tqdm import tqdm

```

```

train_losses = []

```

```

test_losses = []
train_acc = []
test_acc = []

def train(model, device, train_loader, optimizer, epoch):
    model.train()
    pbar = tqdm(train_loader)
    correct = 0
    processed = 0
    for batch_idx, (data, target) in enumerate(pbar):
        # get samples
        data, target = data.to(device), target.to(device)

        # Init
        optimizer.zero_grad()
        # In PyTorch, we need to set the gradients to zero before starting to do backpropagation because PyTorch accumulates
        # Because of this, when you start your training loop, ideally you should zero out the gradients so that you do the par

        # Predict
        y_pred = model(data)

        # Calculate loss
        loss = F.nll_loss(y_pred, target)
        train_losses.append(loss)

        # Backpropagation
        loss.backward()
        optimizer.step()

        # Update pbar-tqdm

        pred = y_pred.argmax(dim=1, keepdim=True) # get the index of the max log-probability
        correct += pred.eq(target.view_as(pred)).sum().item()
        processed += len(data)

    pbar.set_description(desc= f'Loss={loss.item()} Batch_id={batch_idx} Accuracy={100*correct/processed:0.2f}')
    train_acc.append(100*correct/processed)

def test(model, device, test_loader):
    model.eval()
    test_loss = 0
    correct = 0
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            test_loss += F.nll_loss(output, target, reduction='sum').item() # sum up batch loss
            pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-probability
            correct += pred.eq(target.view_as(pred)).sum().item()

    test_loss /= len(test_loader.dataset)
    test_losses.append(test_loss)

    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.2f}%)\n'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))

    test_acc.append(100. * correct / len(test_loader.dataset))

from torch.optim.lr_scheduler import StepLR

model = Net().to(device)
optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
# scheduler = StepLR(optimizer, step_size=6, gamma=0.1)

EPOCHS = 15
for epoch in range(EPOCHS):
    print("EPOCH:", epoch)
    train(model, device, train_loader, optimizer, epoch)
    # scheduler.step()
    test(model, device, test_loader)

    EPOCH: 0
    Loss=0.04363084211945534 Batch_id=468 Accuracy=96.39: 100%|██████████| 469/469 [01:10<00:00, 6.63it/s]

    Test set: Average loss: 0.0355, Accuracy: 9881/10000 (98.81%)

    EPOCH: 1
    Loss=0.014005164615809917 Batch_id=468 Accuracy=98.99: 100%|██████████| 469/469 [01:13<00:00, 6.41it/s]

```

```

Test set: Average loss: 0.0262, Accuracy: 9915/10000 (99.15%)

EPOCH: 2
Loss=0.07649312913417816 Batch_id=468 Accuracy=99.21: 100%|██████████| 469/469 [01:13<00:00, 6.41it/s]

Test set: Average loss: 0.0242, Accuracy: 9915/10000 (99.15%)

EPOCH: 3
Loss=0.001318302471190691 Batch_id=468 Accuracy=99.38: 100%|██████████| 469/469 [01:13<00:00, 6.38it/s]

Test set: Average loss: 0.0309, Accuracy: 9906/10000 (99.06%)

EPOCH: 4
Loss=0.008248790167272091 Batch_id=468 Accuracy=99.51: 100%|██████████| 469/469 [01:13<00:00, 6.38it/s]

Test set: Average loss: 0.0277, Accuracy: 9908/10000 (99.08%)

EPOCH: 5
Loss=0.01677345670759678 Batch_id=468 Accuracy=99.64: 100%|██████████| 469/469 [01:13<00:00, 6.37it/s]

Test set: Average loss: 0.0210, Accuracy: 9930/10000 (99.30%)

EPOCH: 6
Loss=0.01781732775270939 Batch_id=468 Accuracy=99.65: 100%|██████████| 469/469 [01:13<00:00, 6.37it/s]

Test set: Average loss: 0.0211, Accuracy: 9934/10000 (99.34%)

EPOCH: 7
Loss=0.0012531877728179097 Batch_id=468 Accuracy=99.70: 100%|██████████| 469/469 [01:13<00:00, 6.39it/s]

Test set: Average loss: 0.0160, Accuracy: 9949/10000 (99.49%)

EPOCH: 8
Loss=0.0028944441583007574 Batch_id=468 Accuracy=99.74: 100%|██████████| 469/469 [01:13<00:00, 6.38it/s]

Test set: Average loss: 0.0211, Accuracy: 9934/10000 (99.34%)

EPOCH: 9
Loss=0.005427862051874399 Batch_id=468 Accuracy=99.82: 100%|██████████| 469/469 [01:13<00:00, 6.37it/s]

Test set: Average loss: 0.0164, Accuracy: 9949/10000 (99.49%)

EPOCH: 10
Loss=0.0012176345335319638 Batch_id=468 Accuracy=99.87: 100%|██████████| 469/469 [01:13<00:00, 6.37it/s]

Test set: Average loss: 0.0158, Accuracy: 9952/10000 (99.52%)

EPOCH: 11
Loss=0.0016973119927570224 Batch_id=468 Accuracy=99.88: 100%|██████████| 469/469 [01:13<00:00, 6.38it/s]

```

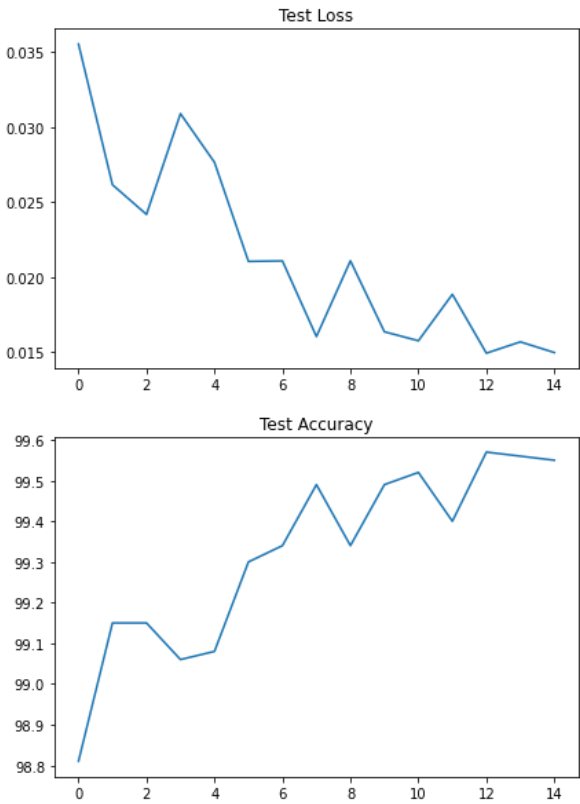
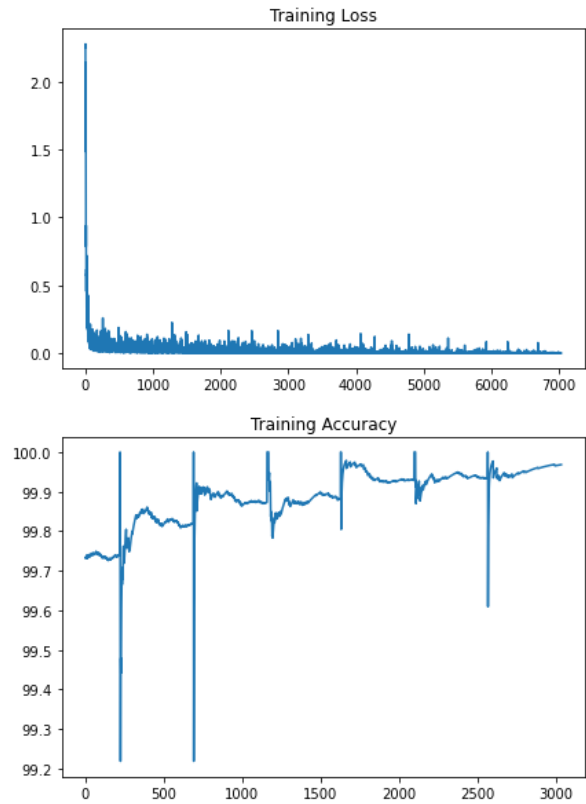
```

%matplotlib inline
import matplotlib.pyplot as plt

train_losses = [i.item() for i in train_losses]
fig, axs = plt.subplots(2,2,figsize=(15,10))
axs[0, 0].plot(train_losses)
axs[0, 0].set_title("Training Loss")
axs[1, 0].plot(train_acc[4000:])
axs[1, 0].set_title("Training Accuracy")
axs[0, 1].plot(test_losses)
axs[0, 1].set_title("Test Loss")
axs[1, 1].plot(test_acc)
axs[1, 1].set_title("Test Accuracy")

```

Text(0.5, 1.0, 'Test Accuracy')



[Colab paid products](#) - [Cancel contracts here](#)

