

Target:

```
-- In the previous step we had 9,970 parameters. In this step we want to
improve this accuracy by using dropouts. We will use a dropout rate of 0.05
```

Results:

```
-- Parameters: 9,970
-- Best Training Accuracy: 98.07
-- Best Test Accuracy: 99.33
```

Analysis:

```
-- There is a slight drop in test accuracy after introducing dropouts.
-- In last few epochs, we can see that the training accuracy is increasing but the test accuracy is going up and down. This might be
```

▼ Import libraries

```
from __future__ import print_function
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
```

▼ Data Transformations (without normalization)

```
# Train Phase transformations
train_transforms = transforms.Compose([
    transforms.ToTensor()
])

# Test Phase transformations
test_transforms = transforms.Compose([
    transforms.ToTensor()
])
```

▼ Dataset and Creating Train/Test Split (without normalization)

```
train = datasets.MNIST('./data', train=True, download=True, transform=train_transforms)
test = datasets.MNIST('./data', train=False, download=True, transform=test_transforms)
```

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>

▼ Dataloader Arguments & Test/Train Dataloaders (without normalization)

```
SEED = 1

# CUDA?
cuda = torch.cuda.is_available()
print("CUDA Available?", cuda)

# For reproducibility
torch.manual_seed(SEED)

if cuda:
    torch.cuda.manual_seed(SEED)

# dataloader arguments - something you'll fetch these from cmdprmt
dataloader_args = dict(shuffle=True, batch_size=128, num_workers=4, pin_memory=True) if cuda else dict(shuffle=True, batch

# train dataloader
train_loader = torch.utils.data.DataLoader(train, **dataloader_args)

# test dataloader
test_loader = torch.utils.data.DataLoader(test, **dataloader_args)

CUDA Available? True
/usr/local/lib/python3.8/dist-packages/torch/utils/data/dataloader.py:554: UserWarning: This DataLoader will create 4
warnings.warn(_create_warning_msg(
```

▼ Getting data statistics (without normalization)

We will use the mean and standard deviation that we get from code below to normalize the data

```
import numpy as np

train_data = train.train_data
train_data = train.transform(train_data.numpy())

print('[Train]')
print(' - Numpy Shape:', train.train_data.cpu().numpy().shape)
print(' - Tensor Shape:', train.train_data.size())
print(' - min:', torch.min(train_data))
print(' - max:', torch.max(train_data))
print(' - mean:', torch.mean(train_data))
print(' - std:', torch.std(train_data))
print(' - var:', torch.var(train_data))

dataiter = iter(train_loader)
images, labels = next(dataiter)

print(images.shape)
print(labels.shape)

# Let's visualize some of the images
%matplotlib inline
import matplotlib.pyplot as plt

plt.imshow(images[0].numpy().squeeze(), cmap='gray_r')
```

```

/usr/local/lib/python3.8/dist-packages/torchvision/datasets/mnist.py:75: Us
warnings.warn("train_data has been renamed data")
[Train]
- Numpy Shape: (60000, 28, 28)
- Tensor Shape: torch.Size([60000, 28, 28])
- min: tensor(0.)
- max: tensor(1.)
- mean: tensor(0.1307)
- std: tensor(0.3081)
- var: tensor(0.0949)
torch.Size([100, 1, 28, 28])

```

▾ Data Transformations (with normalization)

```

# Train Phase transformations
train_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

# Test Phase transformations
test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.1307,), (0.3081,))
])

```

▾ Dataset and Creating Train/Test Split (with normalization)

```

train = datasets.MNIST('./data', train=True, download=True, transform=train_transforms)
test = datasets.MNIST('./data', train=False, download=True, transform=test_transforms)

```

▾ Dataloader Arguments & Test/Train Dataloaders (with normalization)

```

SEED = 1

# CUDA?
cuda = torch.cuda.is_available()
print("CUDA Available?", cuda)

# For reproducibility
torch.manual_seed(SEED)

if cuda:
    torch.cuda.manual_seed(SEED)

# dataloader arguments - something you'll fetch these from cmdprmt
dataloader_args = dict(shuffle=True, batch_size=128, num_workers=4, pin_memory=True) if cuda else dict(shuffle=True, batch

# train dataloader
train_loader = torch.utils.data.DataLoader(train, **dataloader_args)

# test dataloader
test_loader = torch.utils.data.DataLoader(test, **dataloader_args)

CUDA Available? True

```

▾ Getting data statistics (with normalization)

We will use the mean and standard deviation that we get from code below to normalize the data

```

import numpy as np

train_data = train.train_data
train_data = train.transform(train_data.numpy())

print('[Train]')
print(' - Numpy Shape:', train.train_data.cpu().numpy().shape)
print(' - Tensor Shape:', train.train_data.size())
print(' - min:', torch.min(train_data))

```

```

print(' - max:', torch.max(train_data))
print(' - mean:', torch.mean(train_data))
print(' - std:', torch.std(train_data))
print(' - var:', torch.var(train_data))

dataiter = iter(train_loader)
images, labels = next(dataiter)

print(images.shape)
print(labels.shape)

# Let's visualize some of the images
%matplotlib inline
import matplotlib.pyplot as plt


plt.imshow(images[0].numpy().squeeze(), cmap='gray_r')

```

```

[Train]
- Numpy Shape: (60000, 28, 28)
- Tensor Shape: torch.Size([60000, 28, 28])
- min: tensor(-0.4242)
- max: tensor(2.8215)
- mean: tensor(-0.0001)
- std: tensor(1.0000)
- var: tensor(1.0001)
torch.Size([128, 1, 28, 28])
torch.Size([128])
<matplotlib.image.AxesImage at 0x7fc45d805220>

```



Model

```

dropout_value=.05
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()

        #input block
        self.convblock1=nn.Sequential(nn.Conv2d(in_channels=.1, out_channels=.10, kernel_size=.3, padding=.1),
        .....nn.BatchNorm2d(10),
        .....nn.ReLU(),
        .....nn.Dropout(dropout_value)) #R_in=.1, C_in=.28, K=.3, P=.1, S=.1, J_in=.1, J_out=
        ....

        #conv block 1
        self.convblock2=nn.Sequential(nn.Conv2d(in_channels=.10, out_channels=.12, kernel_size=.3, padding=.1),
        .....nn.BatchNorm2d(12),
        .....nn.ReLU(),
        .....nn.Dropout(dropout_value)) #R_in=.3, C_in=.28, K=.3, P=.1, S=.1, J_in=.1, J_out=
        ....

        #conv block 2
        self.convblock3=nn.Sequential(nn.Conv2d(in_channels=.12, out_channels=.14, kernel_size=.3, padding=.1),
        .....nn.BatchNorm2d(14),
        .....nn.ReLU(),
        .....nn.Dropout(dropout_value)) #R_in=.5, C_in=.28, K=.3, P=.1, S=.1, J_in=.1, J_out=
        ....

        #transition block 1
        self.convblock4=nn.Sequential(nn.Conv2d(in_channels=.14, out_channels=.16, kernel_size=.3, padding=.1),
        .....nn.BatchNorm2d(16),
        .....nn.ReLU(),
        .....nn.Dropout(dropout_value)) #R_in=.7, C_in=.28, K=.3, P=.1, S=.1, J_in=.1, J_out=

        self.pool1=nn.MaxPool2d(2, 2) #R_in=.9, C_in=.28, K=.2, P=.0, S=.2, J_in=.1, J_out=.2, R_out=.R_in+.(K-1)*J

        #conv block 3
        self.convblock5=nn.Sequential(nn.Conv2d(in_channels=.16, out_channels=.12, kernel_size=.3, padding=.1),
        .....nn.BatchNorm2d(12),
        .....nn.ReLU(),
        .....nn.Dropout(dropout_value)) #R_in=.10, C_in=.14, K=.3, P=.1, S=.1, J_in=.2, J_out=

```

```

....
....#conv.block*4
....self.convblock6=nn.Sequential(nn.Conv2d(in_channels=12,out_channels=12,kernel_size=3,padding=1),
.....nn.BatchNorm2d(12),
.....nn.ReLU(),
.....nn.Dropout(dropout_value))#R_in=14,C_in=14,K=3,P=1,S=1,J_in=2,J_out

....#conv.block*4
....self.convblock7=nn.Sequential(nn.Conv2d(in_channels=12,out_channels=10,kernel_size=3,padding=1),
.....nn.BatchNorm2d(10),
.....nn.ReLU(),
.....nn.Dropout(dropout_value))#R_in=18,C_in=14,K=3,P=1,S=1,J_in=2,J_out
....#gap.layer*
....self.gap=nn.Sequential(
.....nn.AvgPool2d(kernel_size=4))#R_in=22,C_in=14,K=4,P=1,S=1,J_in=2,J_out=2,R_out=R_in+

....#output.block
....self.convblock8=nn.Sequential(nn.Conv2d(in_channels=10,out_channels=10,kernel_size=3,padding=0),nn.Dropout

..def forward(self,x):
...x=self.convblock1(x)
...x=self.convblock2(x)
...x=self.convblock3(x)
...x=self.convblock4(x)
...x=self.pool1(x)
...x=self.convblock5(x)
...x=self.convblock6(x)
...x=self.convblock7(x)
...x=self.gap(x)
...x=self.convblock8(x)
...x=x.view(-1,10)
...return F.log_softmax(x,dim=-1)

```

Model parameters

```

!pip install torchsummary
from torchsummary import summary

```

```

use_cuda = torch.cuda.is_available()
device = torch.device("cuda" if use_cuda else "cpu")

```

```

model = Net().to(device)
summary(model, input_size = (1, 28, 28))

```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
Requirement already satisfied: torchsummary in /usr/local/lib/python3.8/dist-packages (1.5.1)

Layer (type)	Output Shape	Param #

Conv2d-1	[-1, 10, 28, 28]	100
BatchNorm2d-2	[-1, 10, 28, 28]	20
ReLU-3	[-1, 10, 28, 28]	0
Dropout-4	[-1, 10, 28, 28]	0
Conv2d-5	[-1, 12, 28, 28]	1,092
BatchNorm2d-6	[-1, 12, 28, 28]	24
ReLU-7	[-1, 12, 28, 28]	0
Dropout-8	[-1, 12, 28, 28]	0
Conv2d-9	[-1, 14, 28, 28]	1,526
BatchNorm2d-10	[-1, 14, 28, 28]	28
ReLU-11	[-1, 14, 28, 28]	0
Dropout-12	[-1, 14, 28, 28]	0
Conv2d-13	[-1, 16, 28, 28]	2,032
BatchNorm2d-14	[-1, 16, 28, 28]	32
ReLU-15	[-1, 16, 28, 28]	0
Dropout-16	[-1, 16, 28, 28]	0
MaxPool2d-17	[-1, 16, 14, 14]	0
Conv2d-18	[-1, 12, 14, 14]	1,740
BatchNorm2d-19	[-1, 12, 14, 14]	24
ReLU-20	[-1, 12, 14, 14]	0
Dropout-21	[-1, 12, 14, 14]	0
Conv2d-22	[-1, 12, 14, 14]	1,308
BatchNorm2d-23	[-1, 12, 14, 14]	24
ReLU-24	[-1, 12, 14, 14]	0
Dropout-25	[-1, 12, 14, 14]	0
Conv2d-26	[-1, 10, 14, 14]	1,090
BatchNorm2d-27	[-1, 10, 14, 14]	20
ReLU-28	[-1, 10, 14, 14]	0
Dropout-29	[-1, 10, 14, 14]	0
AvgPool2d-30	[-1, 10, 3, 3]	0
Conv2d-31	[-1, 10, 1, 1]	910
Dropout-32	[-1, 10, 1, 1]	0

```

=====
Total params: 9,970
Trainable params: 9,970
Non-trainable params: 0
-----
Input size (MB): 0.00
Forward/backward pass size (MB): 1.47
Params size (MB): 0.04
Estimated Total Size (MB): 1.51
-----

```

Training and Testing

```

from tqdm import tqdm

train_losses = []
test_losses = []
train_acc = []
test_acc = []

def train(model, device, train_loader, optimizer, epoch):
    model.train()
    pbar = tqdm(train_loader)
    correct = 0
    processed = 0
    for batch_idx, (data, target) in enumerate(pbar):
        # get samples
        data, target = data.to(device), target.to(device)

        # Init
        optimizer.zero_grad()
        # In PyTorch, we need to set the gradients to zero before starting to do backpropagation because PyTorch accumulates
        # Because of this, when you start your training loop, ideally you should zero out the gradients so that you do the par

        # Predict
        y_pred = model(data)

        # Calculate loss
        loss = F.nll_loss(y_pred, target)
        train_losses.append(loss)

        # Backpropagation
        loss.backward()
        optimizer.step()

        # Update pbar-tqdm

        pred = y_pred.argmax(dim=1, keepdim=True) # get the index of the max log-probability
        correct += pred.eq(target.view_as(pred)).sum().item()
        processed += len(data)

        pbar.set_description(desc= f'Loss={loss.item()} Batch_id={batch_idx} Accuracy={100*correct/processed:0.2f}')
        train_acc.append(100*correct/processed)

def test(model, device, test_loader):
    model.eval()
    test_loss = 0
    correct = 0
    with torch.no_grad():
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            test_loss += F.nll_loss(output, target, reduction='sum').item() # sum up batch loss
            pred = output.argmax(dim=1, keepdim=True) # get the index of the max log-probability
            correct += pred.eq(target.view_as(pred)).sum().item()

    test_loss /= len(test_loader.dataset)
    test_losses.append(test_loss)

    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{} ({:.2f}%)\n'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))

    test_acc.append(100. * correct / len(test_loader.dataset))

from torch.optim.lr_scheduler import StepLR

model = Net().to(device)

```

```

optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
# scheduler = StepLR(optimizer, step_size=6, gamma=0.1)

EPOCHS = 15
for epoch in range(EPOCHS):
    print("EPOCH:", epoch)
    train(model, device, train_loader, optimizer, epoch)
    # scheduler.step()
    test(model, device, test_loader)

    EPOCH: 0
    Loss=0.09847403317689896 Batch_id=468 Accuracy=88.59: 100%|██████████| 469/469 [00:19<00:00, 24.44it/s]

    Test set: Average loss: 0.0915, Accuracy: 9720/10000 (97.20%)

    EPOCH: 1
    Loss=0.05309741571545601 Batch_id=468 Accuracy=96.39: 100%|██████████| 469/469 [00:20<00:00, 23.25it/s]

    Test set: Average loss: 0.0486, Accuracy: 9851/10000 (98.51%)

    EPOCH: 2
    Loss=0.018912313506007195 Batch_id=468 Accuracy=97.04: 100%|██████████| 469/469 [00:15<00:00, 30.16it/s]

    Test set: Average loss: 0.0480, Accuracy: 9848/10000 (98.48%)

    EPOCH: 3
    Loss=0.022453108802437782 Batch_id=468 Accuracy=97.13: 100%|██████████| 469/469 [00:16<00:00, 27.77it/s]

    Test set: Average loss: 0.0407, Accuracy: 9872/10000 (98.72%)

    EPOCH: 4
    Loss=0.039117444306612015 Batch_id=468 Accuracy=97.51: 100%|██████████| 469/469 [00:15<00:00, 29.87it/s]

    Test set: Average loss: 0.0399, Accuracy: 9872/10000 (98.72%)

    EPOCH: 5
    Loss=0.020529208704829216 Batch_id=468 Accuracy=97.68: 100%|██████████| 469/469 [00:15<00:00, 29.99it/s]

    Test set: Average loss: 0.0324, Accuracy: 9900/10000 (99.00%)

    EPOCH: 6
    Loss=0.08593115955591202 Batch_id=468 Accuracy=97.67: 100%|██████████| 469/469 [00:16<00:00, 28.35it/s]

    Test set: Average loss: 0.0297, Accuracy: 9912/10000 (99.12%)

    EPOCH: 7
    Loss=0.035492051392793655 Batch_id=468 Accuracy=97.70: 100%|██████████| 469/469 [00:15<00:00, 30.35it/s]

    Test set: Average loss: 0.0370, Accuracy: 9878/10000 (98.78%)

    EPOCH: 8
    Loss=0.0680157020688057 Batch_id=468 Accuracy=97.77: 100%|██████████| 469/469 [00:15<00:00, 30.58it/s]

    Test set: Average loss: 0.0235, Accuracy: 9927/10000 (99.27%)

    EPOCH: 9
    Loss=0.0273868886165619 Batch_id=468 Accuracy=97.84: 100%|██████████| 469/469 [00:15<00:00, 30.25it/s]

    Test set: Average loss: 0.0246, Accuracy: 9925/10000 (99.25%)

    EPOCH: 10
    Loss=0.03704260662198067 Batch_id=468 Accuracy=97.94: 100%|██████████| 469/469 [00:15<00:00, 29.90it/s]





    Test set: Average loss: 0.0246, Accuracy: 9918/10000 (99.18%)

    EPOCH: 11
    Loss=0.035303425043821335 Batch_id=468 Accuracy=97.97: 100%|██████████| 469/469 [00:16<00:00, 27.68it/s]

train_losses = [i.item() for i in train_losses]
%matplotlib inline
import matplotlib.pyplot as plt
fig, axs = plt.subplots(2,2,figsize=(15,10))
axs[0, 0].plot(train_losses)
axs[0, 0].set_title("Training Loss")
axs[1, 0].plot(train_acc[4000:])
axs[1, 0].set_title("Training Accuracy")
axs[0, 1].plot(test_losses)
axs[0, 1].set_title("Test Loss")
axs[1, 1].plot(test_acc)
axs[1, 1].set_title("Test Accuracy")

```

Text(0.5, 1.0, 'Test Accuracy')



Colab paid products - [Cancel contracts here](#)